

# Implementing DeepSDF

Final Project, IFT6113 - Geometric Modeling and Shape Analysis

Sylvain Laporte - Fall 2020

# Summary

- What is DeepSDF and why is it useful?
- Our implementation
- Results
- An idea of our extension

# What is DeepSDF?

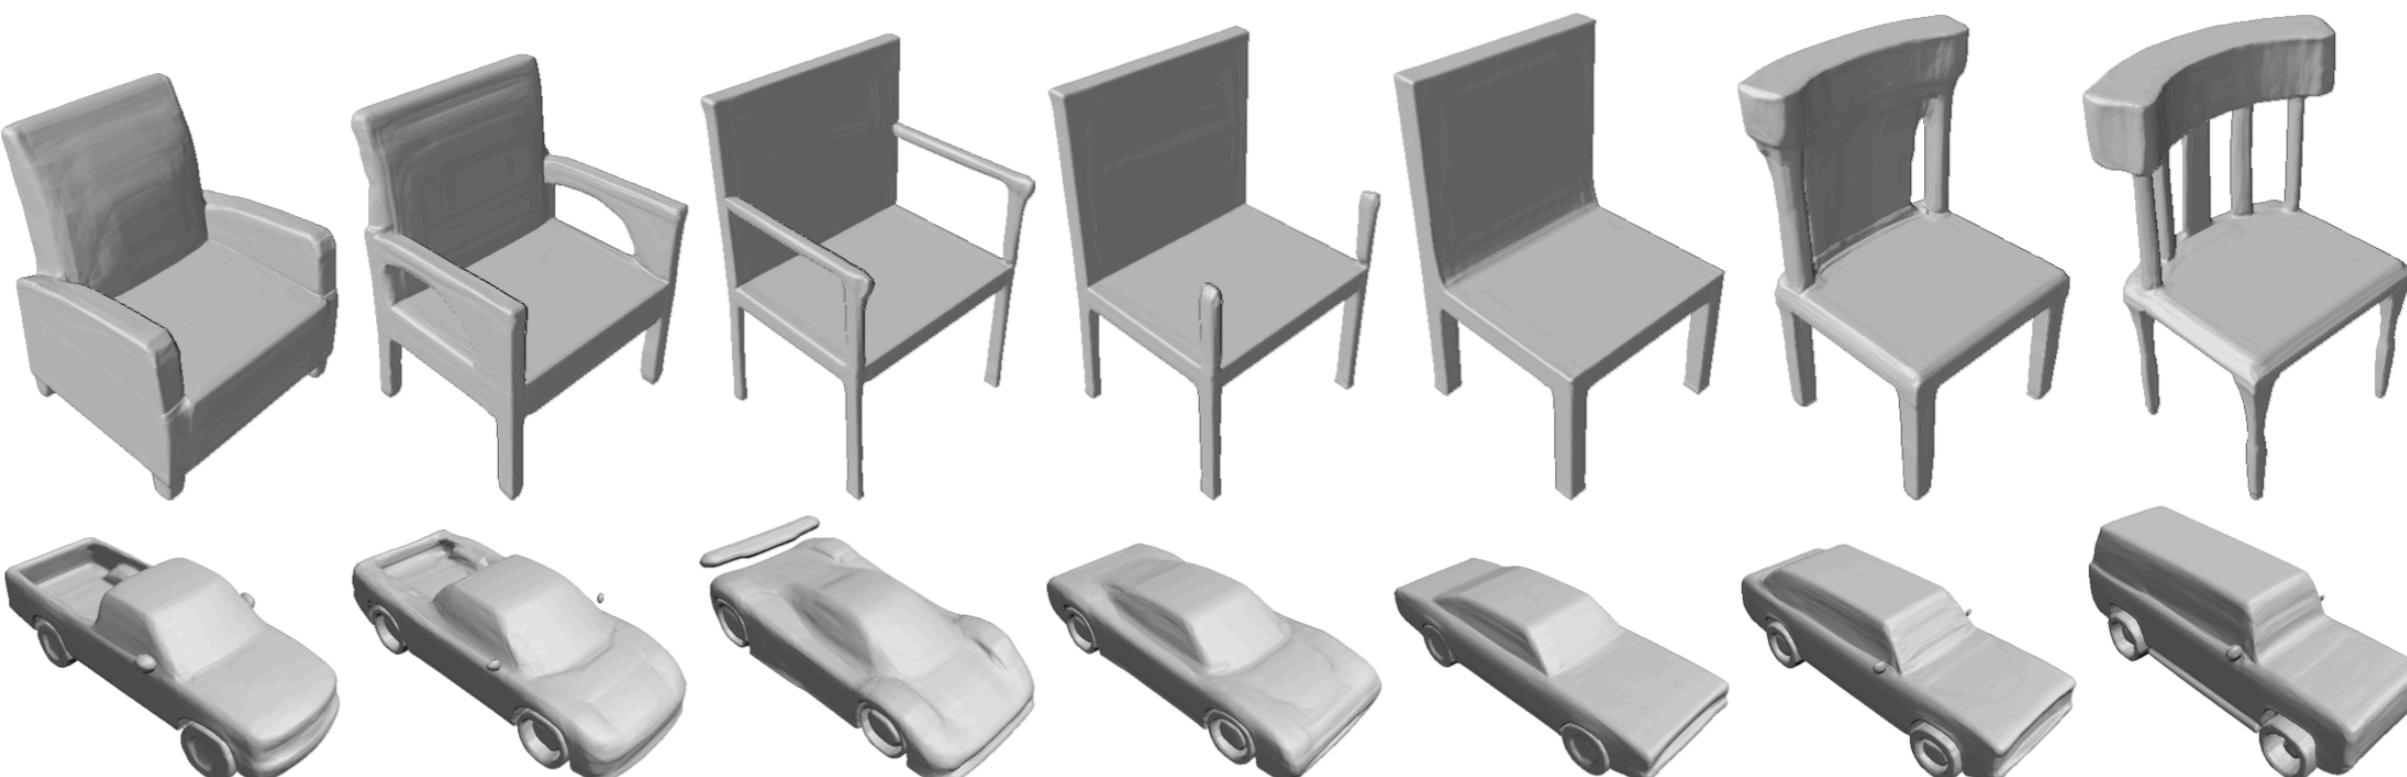
# The paper

## « DeepSDF: Learning Continuous Signed Distances Functions for Shape Representation »

DeepSDF: Learning Continuous Signed Distance Functions for Shape Representation

Jeong Joon Park<sup>1,3†</sup> Peter Florence<sup>2,3†</sup> Julian Straub<sup>3</sup> Richard Newcombe<sup>3</sup> Steven Lovegrove<sup>3</sup>

<sup>1</sup>University of Washington <sup>2</sup>Massachusetts Institute of Technology <sup>3</sup>Facebook Reality Labs



**Figure 1:** DeepSDF represents signed distance functions (SDFs) of shapes via latent code-conditioned feed-forward decoder networks. Above images are raycast renderings of DeepSDF interpolating between two shapes in the learned shape latent space. Best viewed digitally.

**Abstract**

*Computer graphics, 3D computer vision and robotics communities have produced multiple approaches to representing 3D geometry for rendering and reconstruction.*

**1. Introduction**

Deep convolutional networks which are a mainstay of image-based approaches grow quickly in space and time complexity when directly generalized to the 3rd spatial di-

# The problem

**How to represent 3D geometry efficiently for rendering and reconstruction?**

- Computer graphics, 3D computer vision and robotics need a good way to represent 3D geometry:
  - quality
  - flexibility
  - fidelity
  - efficiency in time and memory

# Previous work

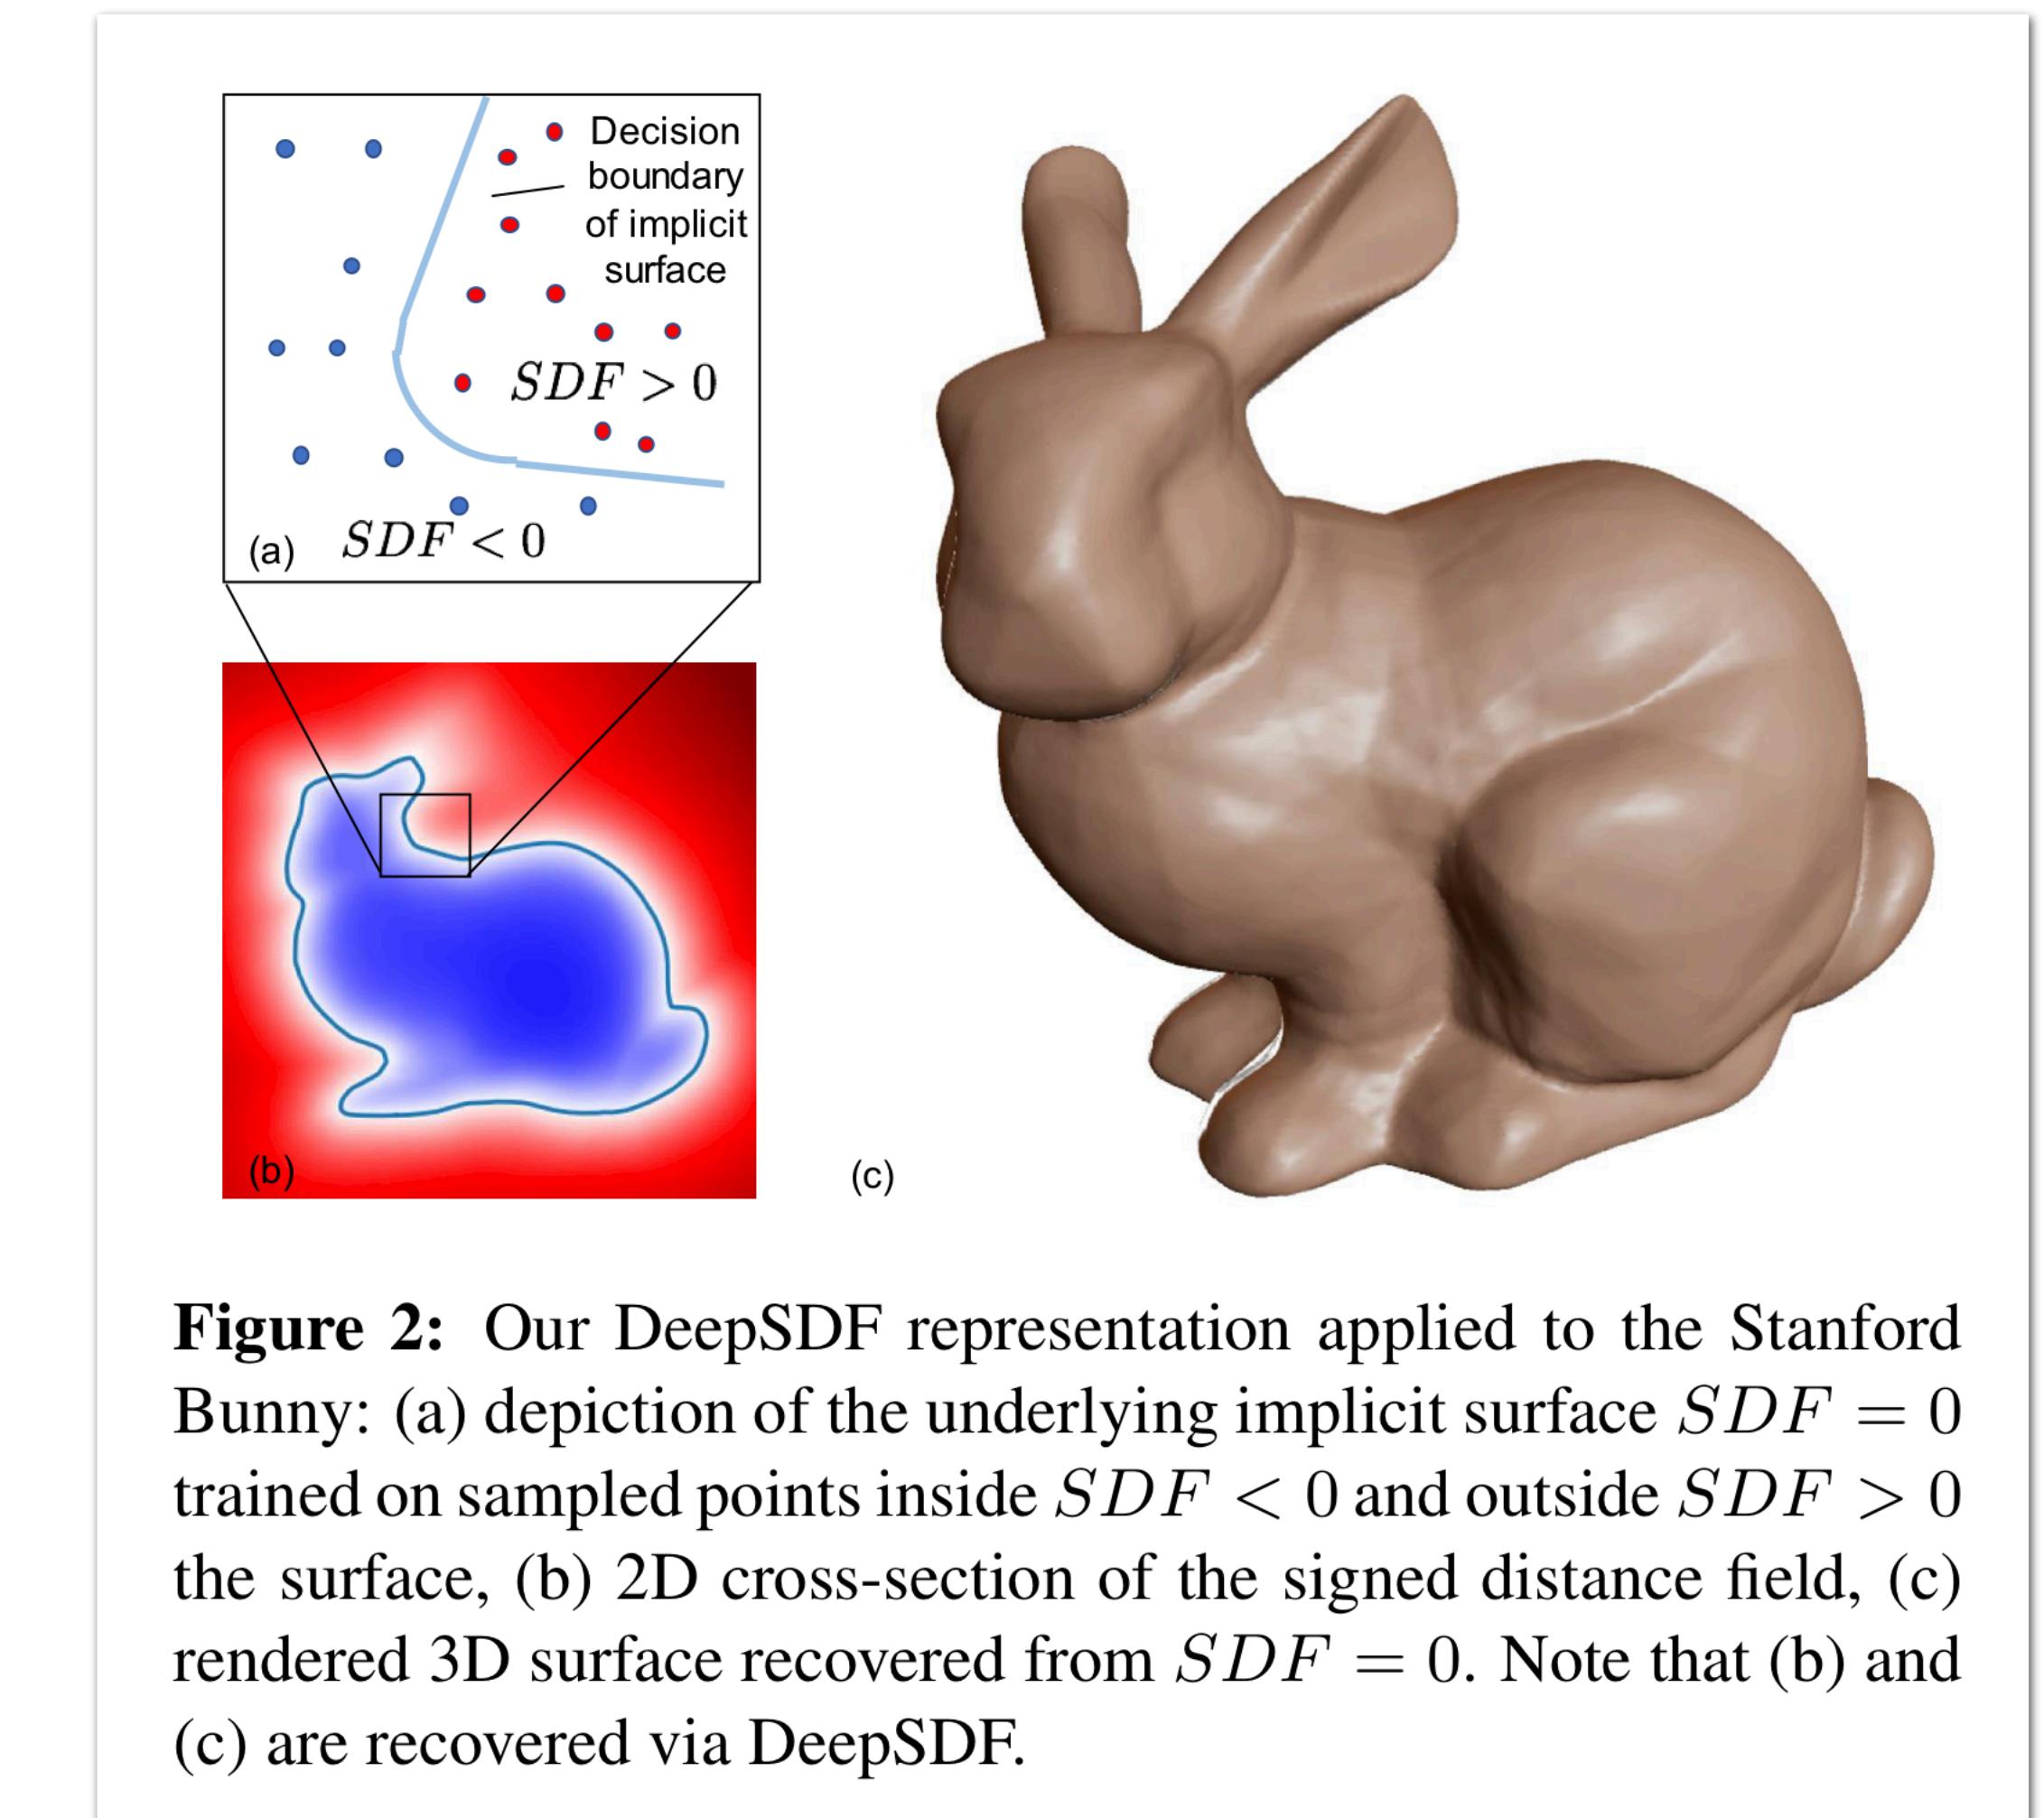
## Problems with existing representations

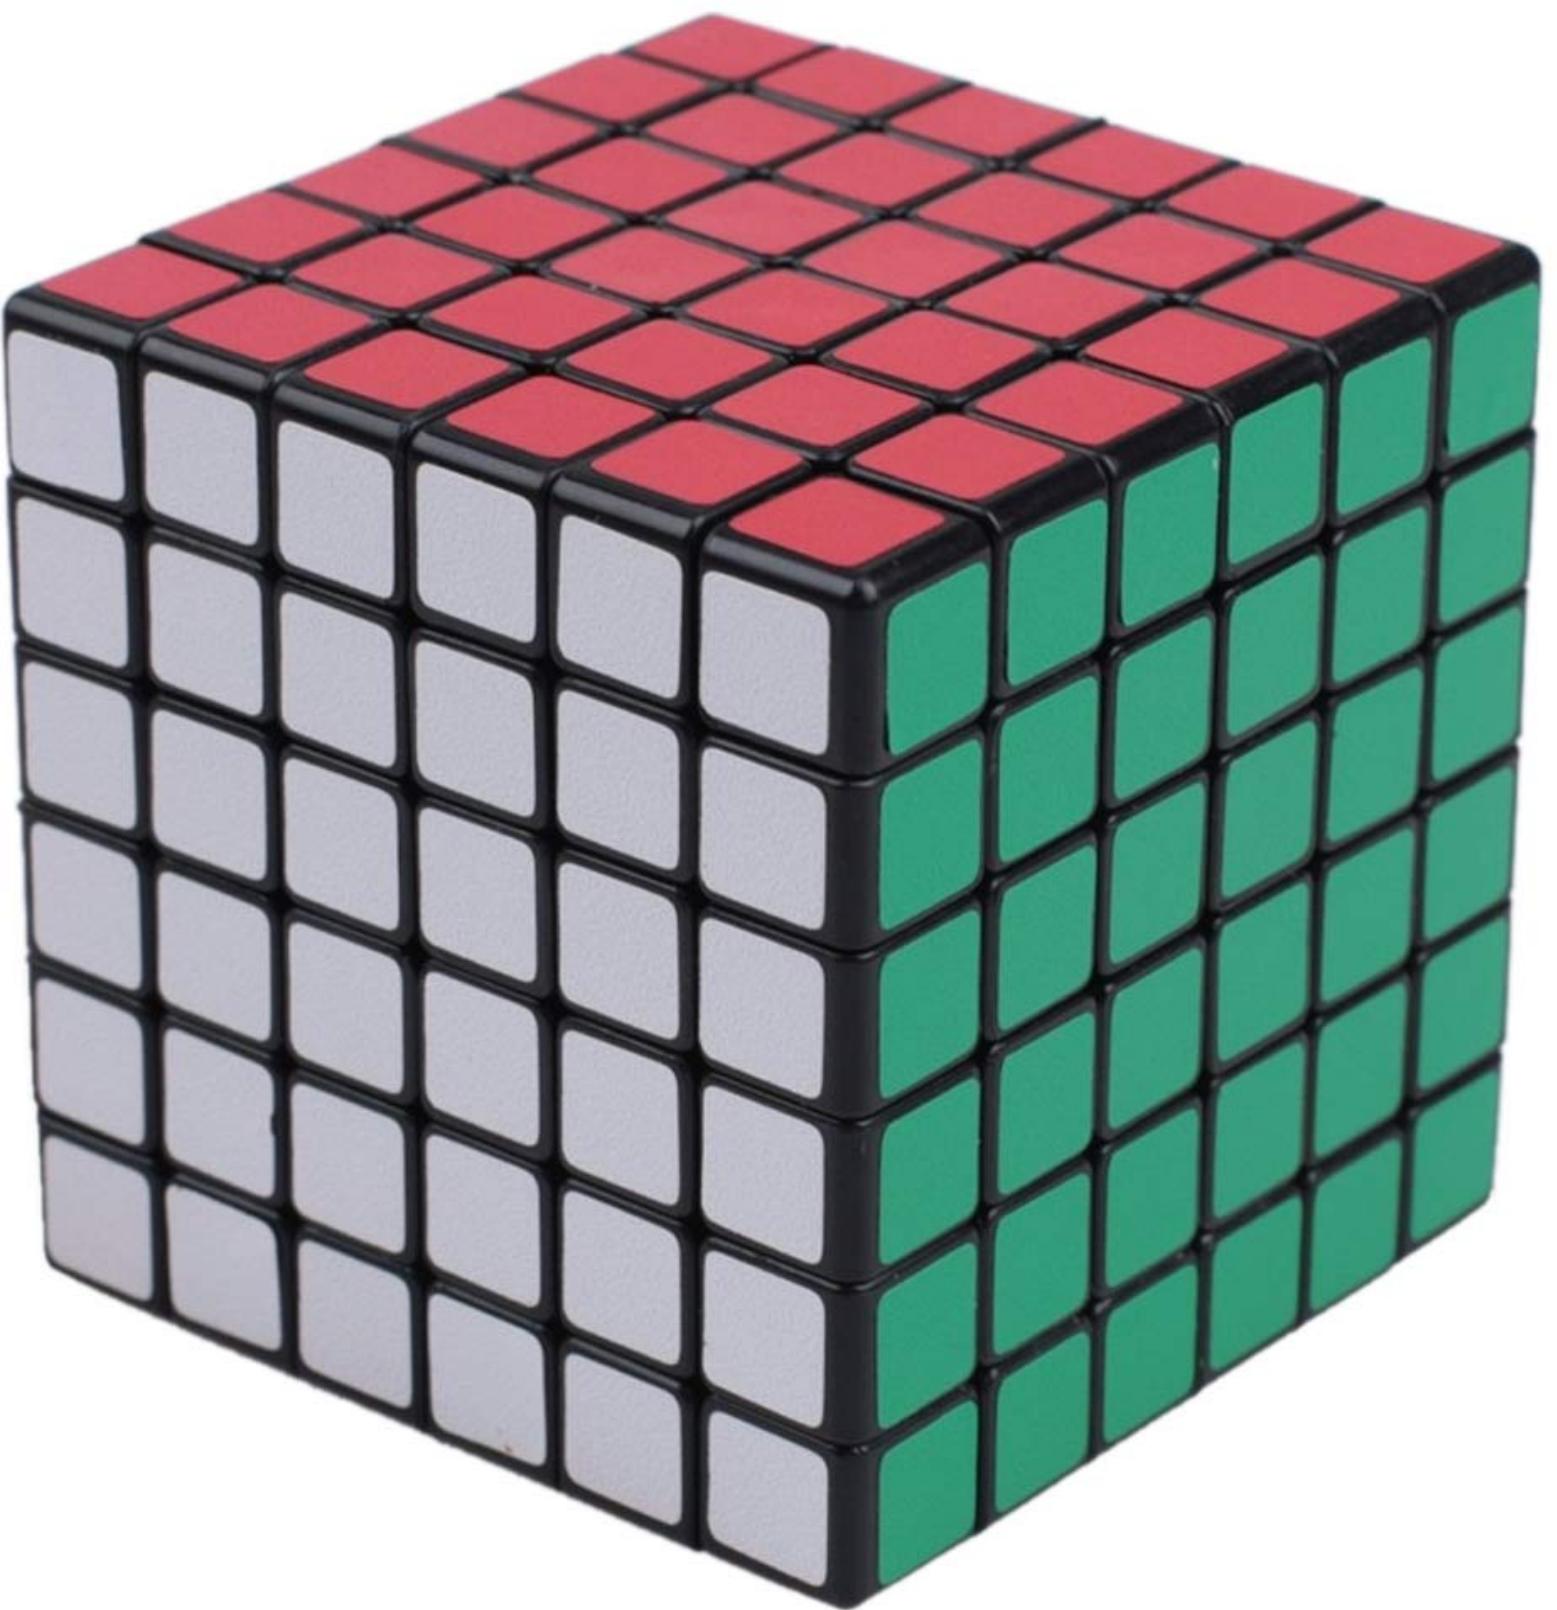
- Deep convolutional networks:
  - grow quickly in space and time complexity when generalized to the 3rd spatial dimension.
- Classical surface representations (triangle or quad meshes):
  - have an unknown number of vertices and an arbitrary topology; hence they are not well suited for training.

# The solution

## Implicit representation using SDF

- $SDF$  = signed distance function
- A continuous function mapping each point in 3D space to its distance to the shape surface.
  - points inside the shape have **negative SDF values**
  - points outside the shape have **positive SDF values**
- SDF is represented as a volumetric field.
- The implicit surface is recovered from  $SDF = 0$ .



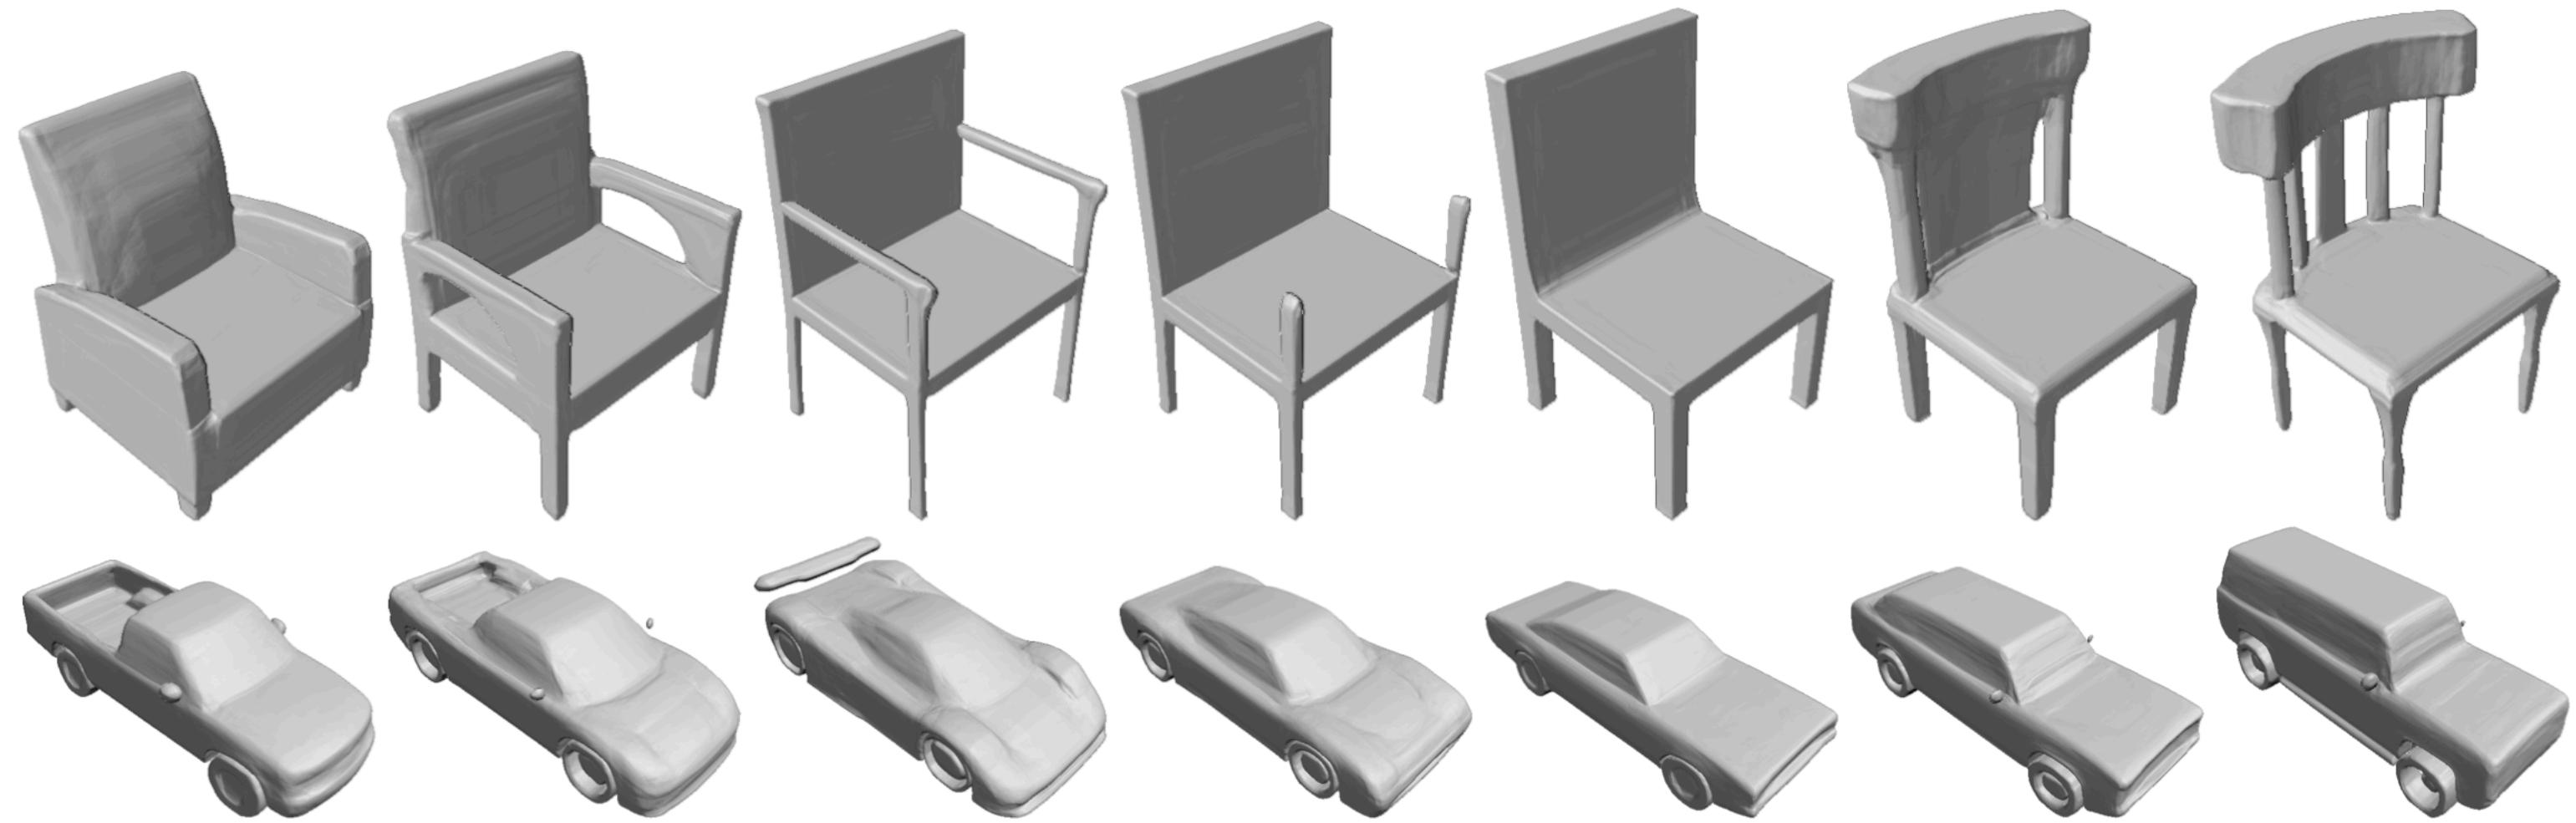


**A volumetric field, made of voxels**

# The solution

## DeepSDF - Learning a 3D generative model of SDFs

- Uses a deep neural network to learn the continuous SDF of a shape from spatial sample points and their known SDF value.
- Then, generate the mesh of that shape from the resulting SDF volumetric field with some technique such as Marching Cubes.
- Advantages:
  - Discretization happens only when the mesh needs to be generated.
  - Can be used to represent entire classes of shapes.

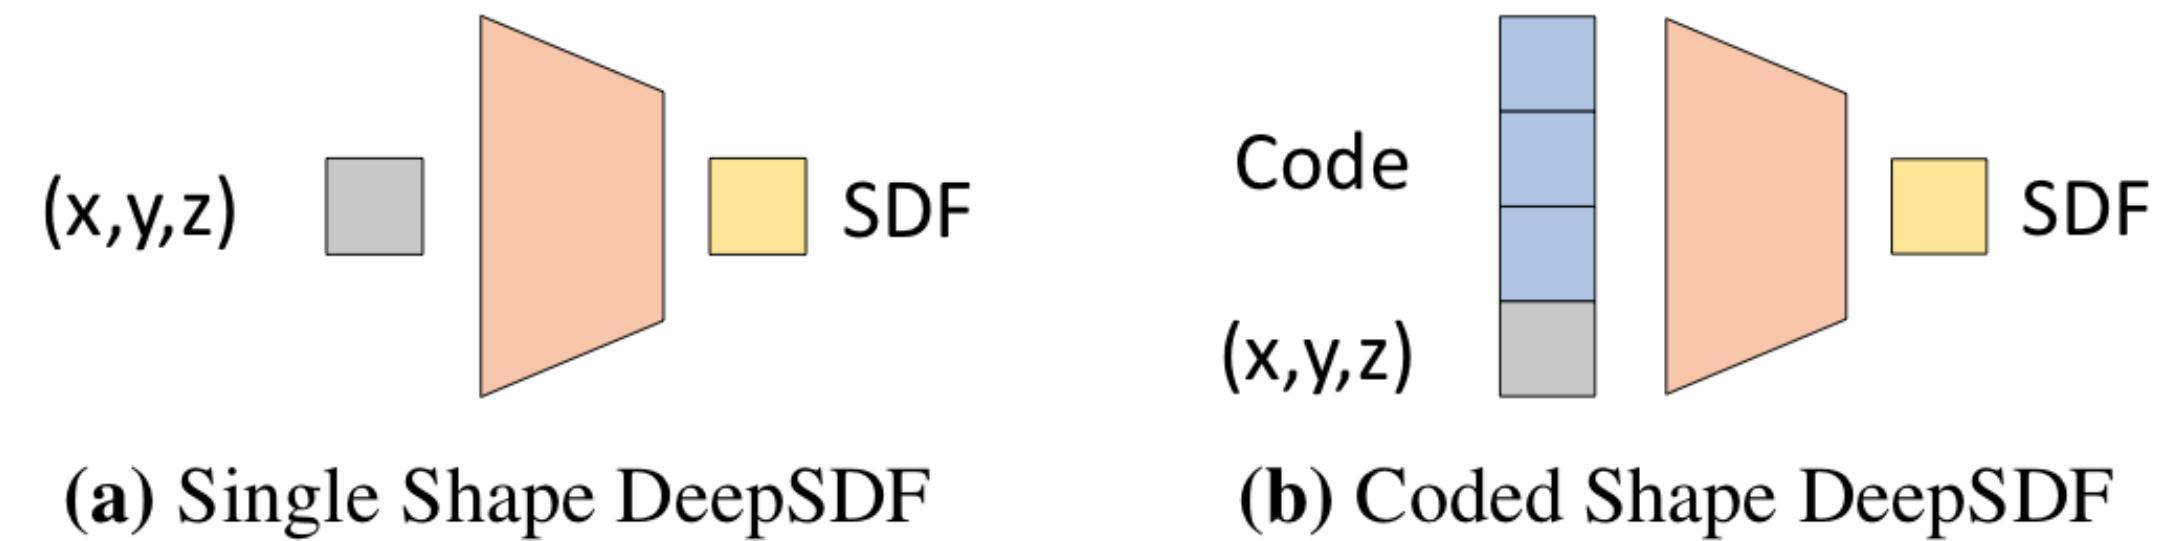


Classes of shapes

# The solution

## Two flavours of DeepSDF

- Single-shape DeepSDF:
  - Learns for one shape.
- Coded-shape DeepSDF:
  - Learns for a class of shapes.
  - Uses an auto-decoder only network.



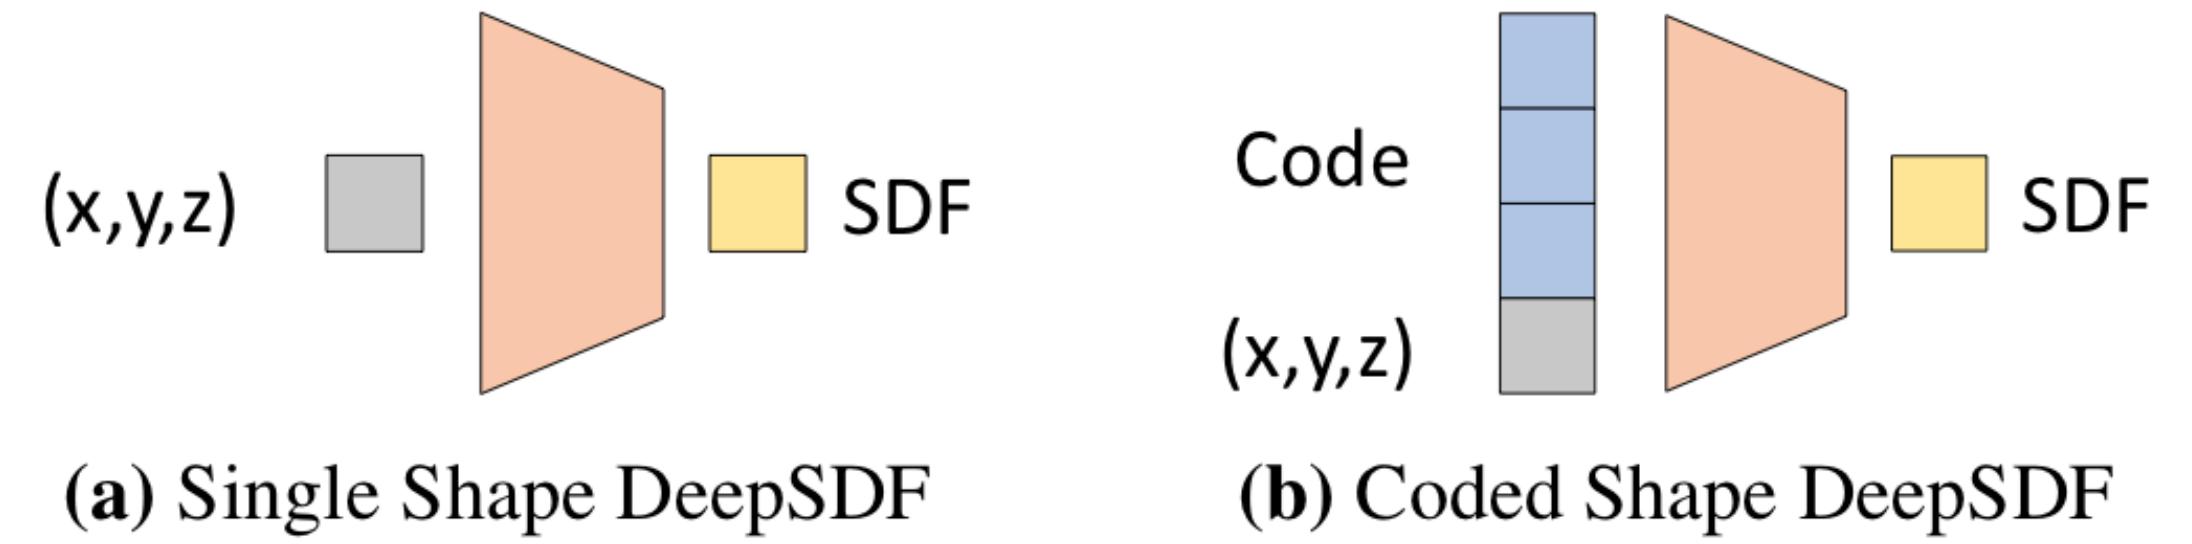
**Figure 3:** In the single-shape DeepSDF instantiation, the shape information is contained in the network itself whereas the coded-shape DeepSDF, the shape information is contained in a code vector that is concatenated with the 3D sample location. In both cases, DeepSDF produces the SDF value at the 3D query location,

# The solution

## Two flavours of DeepSDF

- Single-shape DeepSDF:  
• Learns for one shape.
- Coded-shape DeepSDF:
  - Learns for a class of shapes.
  - Uses an auto-decoder only network.

*Our focus*



**Figure 3:** In the single-shape DeepSDF instantiation, the shape information is contained in the network itself whereas the coded-shape DeepSDF, the shape information is contained in a code vector that is concatenated with the 3D sample location. In both cases, DeepSDF produces the SDF value at the 3D query location,

# Our implementation

# The steps

Let the journey begin...

- A. Prepare the data
- B. Build and train the neural network
- C. Generate a mesh

# The tools

## What we used to code this project

- Python
- PyTorch – for building and training the neural network
- Google Colab – because Apple 😠 CUDA
- trimesh – for importing meshes
- skimage – for marching cubes
- third-party code (mesh-to-sdf) for data preprocessing – because mesh 😠 my code

# Preparing the data

## Which data?

- The network needs a list of tuples, each tuple containing:
  - a spatial sample point
  - its associated true computed SDF value

$$X_i = \{(\mathbf{x}_j, s_j) : s_j = SDF^i(\mathbf{x}_j)\}. \quad (6)$$

# Preparing the data

## How to prepare?

- We are provided complete 3D shape meshes.
- First, we normalize each mesh to a unit sphere.
  - The farthest point must be at a distance  $\leq 1$  from the mesh's center.
  - In practice, we add a small margin, making the sphere's radius =  $1/1.03$ .
- Then, we sample around 250,000 points randomly and compute their SDF.
  - We sample more aggressively near the shape's surface.

# Build and train the neural network

## Neural network architecture

- We train a multi-layer fully-connected feed-forward neural network:
  - 8 fully-connected layers, each applied with dropout.
  - Internal layers are 512-dimensional and have ReLU non-linearities.
  - Output layer has tanh non-linearities.
  - Instead of batch-normalization, we use weight-normalization.
- For single-shape DeepSDF, the network takes a spatial sample point as input.
- For coded-shape DeepSDF, we also add a latent vector as input which represents the desired shape from the specific class to train for.

## The DeepSDF network

```
> ▶ M4
class DeepSDF_single(nn.Module):
    """Multi-layer fully-connected feed-forward neural network with 8 layers, each applied with dropout"""
    def __init__(self, in_size, hidden_size, out_size):
        super().__init__()
        # Hidden layer 1
        self.linear1 = nn.Linear(in_size, hidden_size)
        # Hidden layer 2
        self.linear2 = nn.Linear(hidden_size, hidden_size)
        # Hidden layer 3
        self.linear3 = nn.Linear(hidden_size, hidden_size)
        # Hidden layer 4
        self.linear4 = nn.Linear(hidden_size, hidden_size)
        # Hidden layer 5
        self.linear5 = nn.Linear(hidden_size, hidden_size)
        # Hidden layer 6
        self.linear6 = nn.Linear(hidden_size, hidden_size)
        # Hidden layer 7
        self.linear7 = nn.Linear(hidden_size, hidden_size)
        # Output layer
        self.linear8 = nn.Linear(hidden_size, out_size)

    def forward(self, xb):
        out = self.linear1(xb)
        out = F.relu(out)
        out = self.linear2(out)
        out = F.relu(out)
        out = self.linear3(out)
        out = F.relu(out)
        out = self.linear4(out)
        out = F.relu(out)
        out = self.linear5(out)
        out = F.relu(out)
        out = self.linear6(out)
        out = F.relu(out)
        out = self.linear7(out)
        out = F.relu(out)
        out = self.linear8(out)
        return out

    def training_step(self, batch):
        point_samples, sdf_values = batch
        out = self(point_samples)      # generate predictions
        loss = loss_function(out, sdf_values)    # compute loss
        return loss

    def validation_step(self, batch):
        point_samples, sdf_values = batch
        out = self(point_samples)      # generate predictions
        loss = loss_function(out, sdf_values)    # compute loss
        acc = accuracy(out, sdf_values)    # compute accuracy
        return {'val_loss': loss, 'val_acc': acc}

    def validation_epoch_end(self, outputs):
        batch_losses = [x['val_loss'] for x in outputs]
        epoch_loss = torch.stack(batch_losses).mean()    # combine losses
        return {'val_loss': epoch_loss.item()}

        # batch_accs = [x['val_acc'] for x in outputs]
        # epoch_acc = torch.stack(batch_accs).mean()      # combine accuracies
        # return {'val_loss': epoch_loss.item(), 'val_acc': epoch_acc.item()}

    def epoch_end(self, epoch, result):
        print("Epoch [{}], val_loss: {:.4f}".format(epoch, result['val_loss']))
        # print("Epoch [{}], val_loss: {:.4f}, val_acc: {:.4f}".format(epoch, result['val_loss'], result['val_acc']))
```

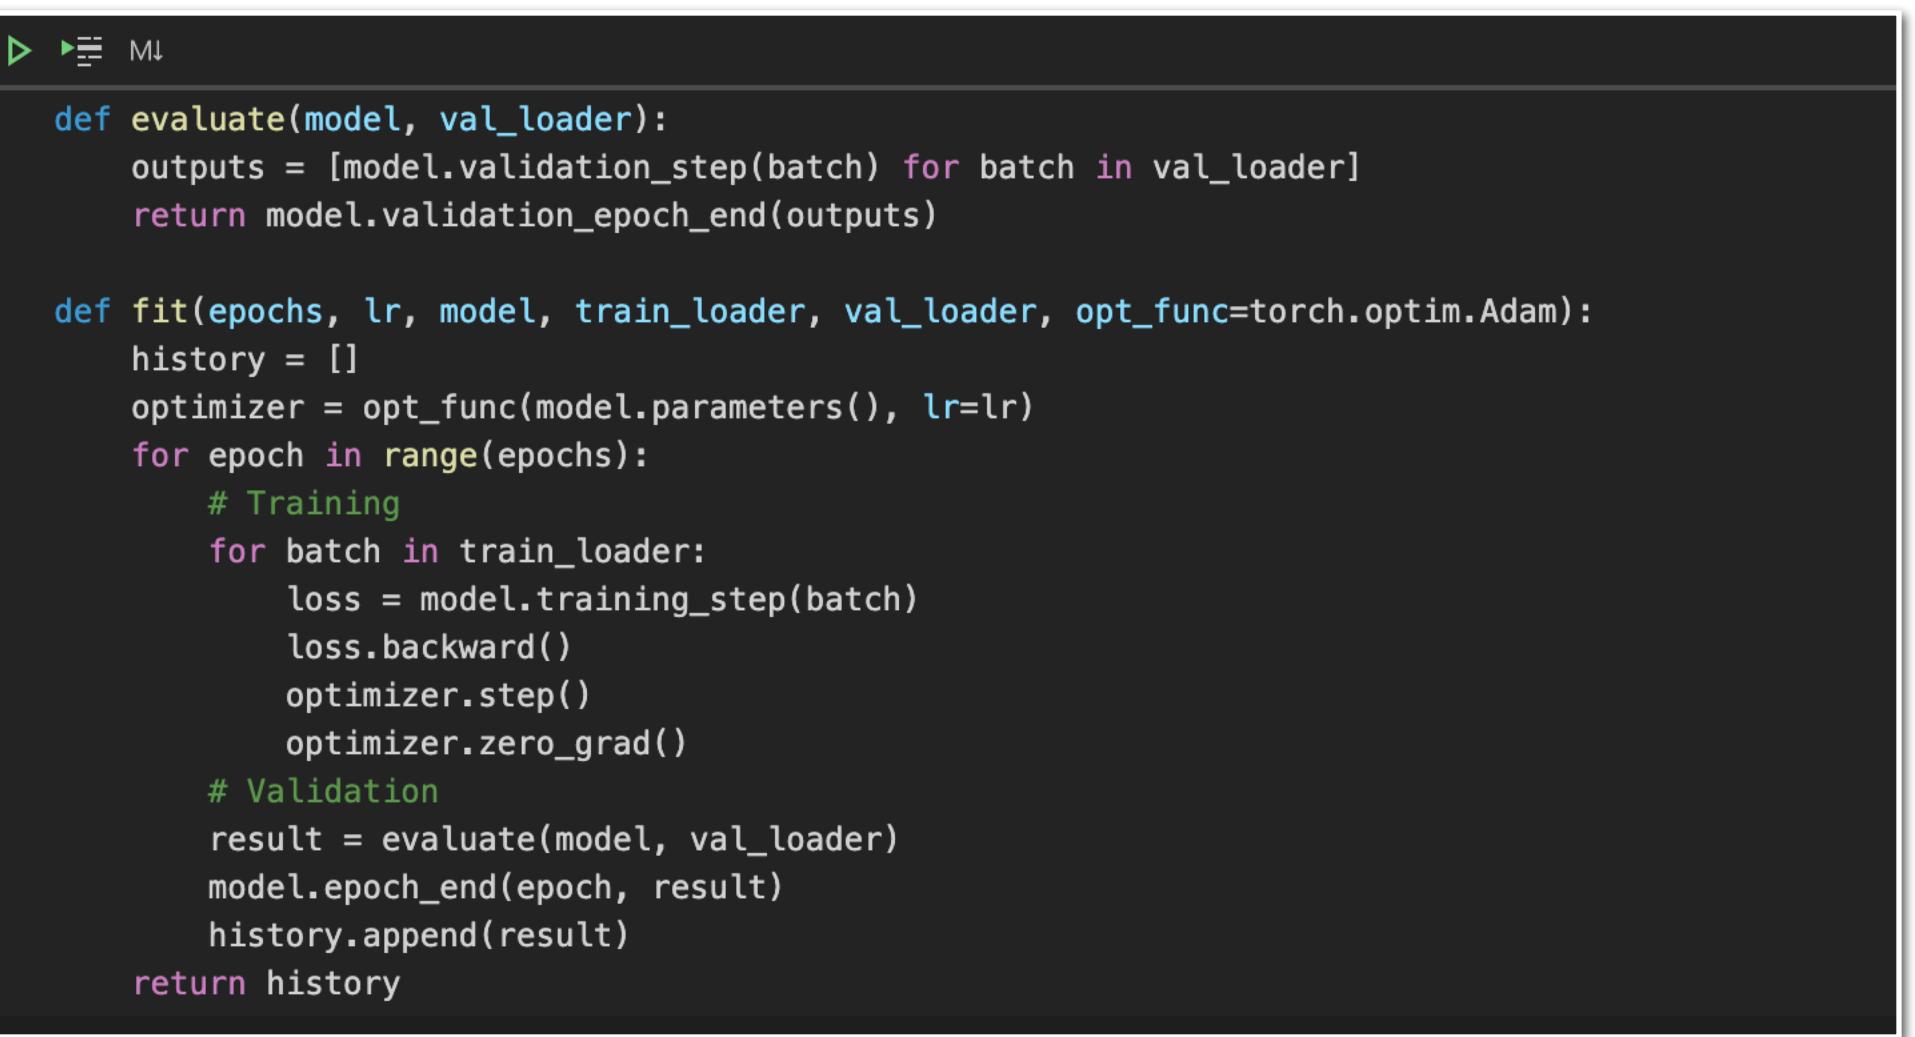
# Building and training the neural network

## Training the network

- We use an L1 loss function.

$$\mathcal{L}(f_\theta(\mathbf{x}), s) = |\text{clamp}(f_\theta(\mathbf{x}), \delta) - \text{clamp}(s, \delta)|, \quad (4)$$

- where  $\text{clamp}(x, \delta) := \min(\delta, \max(-\delta, x))$
- We train the network for 1,000 epochs.
- We use the Adam optimizer.
- We try and tune a learning rate.



```
▶ ▶ Ml
def evaluate(model, val_loader):
    outputs = [model.validation_step(batch) for batch in val_loader]
    return model.validation_epoch_end(outputs)

def fit(epochs, lr, model, train_loader, val_loader, opt_func=torch.optim.Adam):
    history = []
    optimizer = opt_func(model.parameters(), lr=lr)
    for epoch in range(epochs):
        # Training
        for batch in train_loader:
            loss = model.training_step(batch)
            loss.backward()
            optimizer.step()
            optimizer.zero_grad()
        # Validation
        result = evaluate(model, val_loader)
        model.epoch_end(epoch, result)
        history.append(result)
    return history
```

Functions to implement the training

# Generating a mesh

## Voxels and marching cubes

- The trained model approximates the SDF for a specific mesh or a class of meshes.
- We can generate the volumetric field by providing the model with a  $N \times N \times N$  grid of points to evaluate and return their SDF values.
  - $N$  is the resolution of the grid, we used  $N = 100$ .
- We can then give the volumetric field as input to a marching cube algorithm which will compute vertices, faces and normals based on the voxels where  $SDF = 0$ .

```
1 # author: Sylvain Laporte
2 # program: reconstruct_mesh_from_model.py
3 # date: 2020-12-08
4 # object: Reconstructs a mesh from the trained DeepSDF model.
5
6 import torch
7 import deepSDF_model
8 import trimesh
9 from skimage import measure
10
11 # Load the model
12
13 PATH = "trained_models/horse-1-model-10.pt"
14
15 input_size = 3
16 hidden_size = 512
17 out_size = 1
18
19 model = deepSDF_model.DeepSDF_single(input_size, hidden_size=hidden_size, out_size=out_size)
20 model.load_state_dict(torch.load(PATH, map_location=torch.device('cpu')))
21 model.eval()
22
23 # Generate voxels
24
25 RESOLUTION = 100
26 voxels = torch.zeros(RESOLUTION, RESOLUTION, RESOLUTION)
27
28 intervals = [x / RESOLUTION for x in range(RESOLUTION)]
29 points = [[x, y, z] for x in intervals for y in intervals for z in intervals]
30 points = torch.tensor(points)
31
32 sdfs = model(points)
33
34 voxels = sdfs.view(RESOLUTION, RESOLUTION, RESOLUTION).detach().numpy()
35
36 # Generate mesh using marching cubes
37
38 vertices, faces, normals, _ = measure.marching_cubes(voxels, 0)
39 mesh = trimesh.Trimesh(vertices, faces, normals)
40 trimesh.repair.fix_normals(mesh)
41
42 # Show the resulting mesh
43
44 mesh.show()
```

Mesh generating script

# Hurdles

## Lots of new stuff

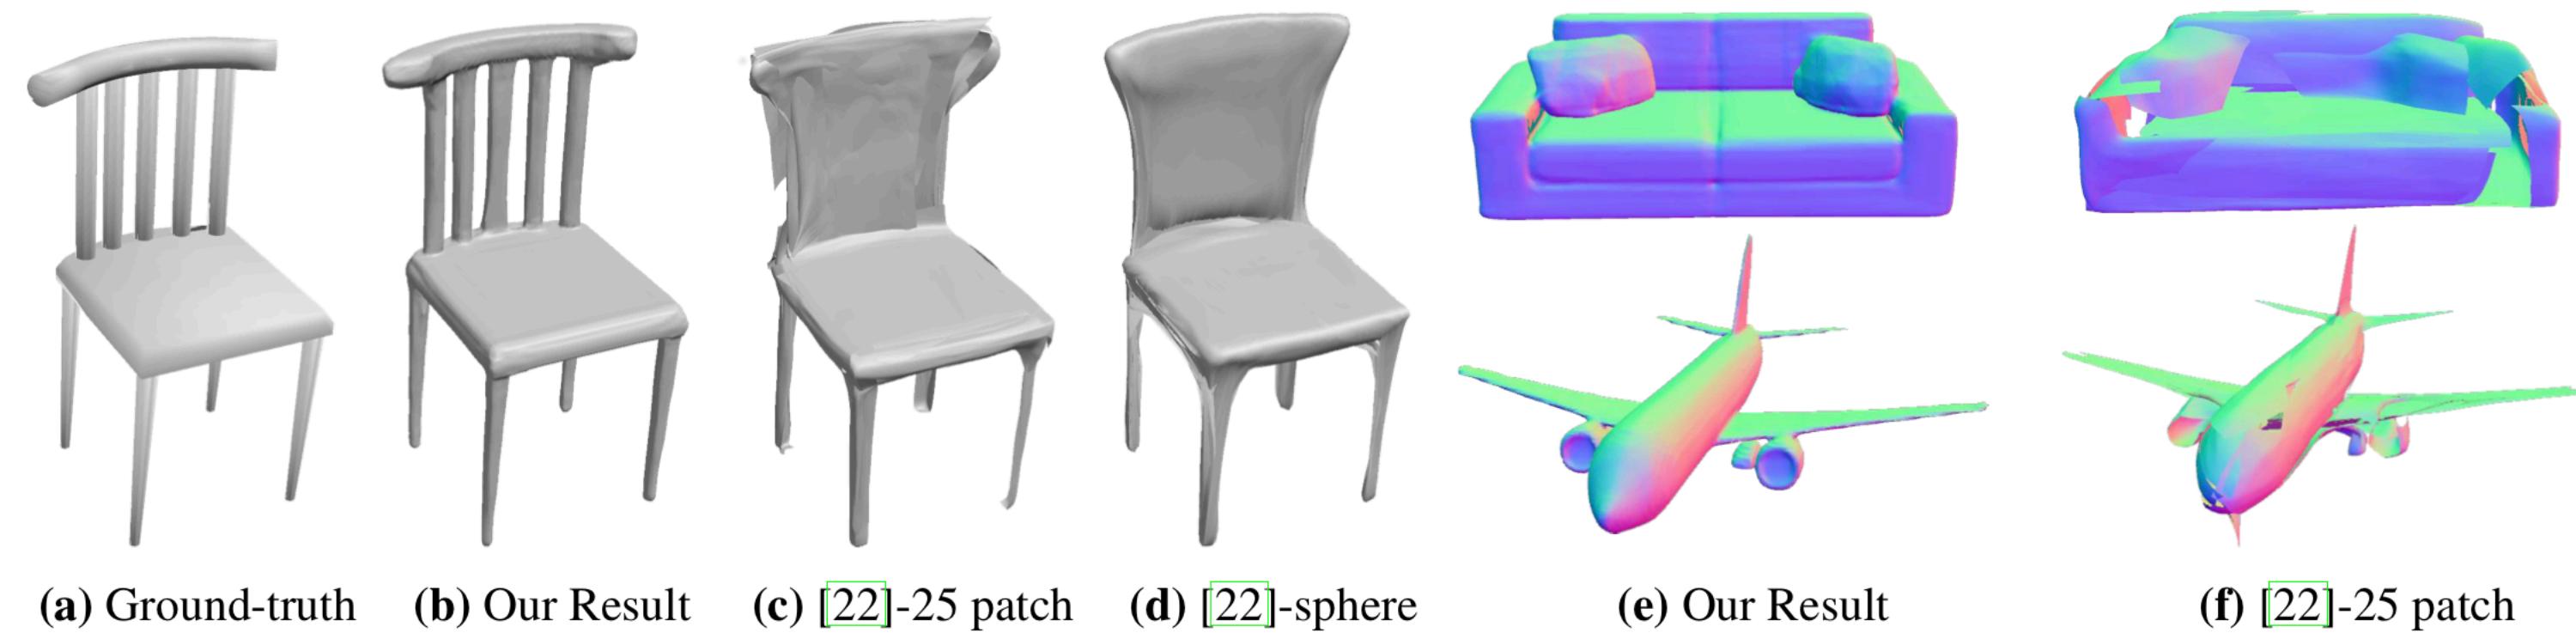
- Getting to know PyTorch beyond MNIST
- Jupyter notebook « eating » my work
- Also, getting to know Google Colab beyond MNIST
  - render context, custom modules, getting things onto the GPU
- What kind of voxels does a marching cubes eats?



# Results

# Expected results

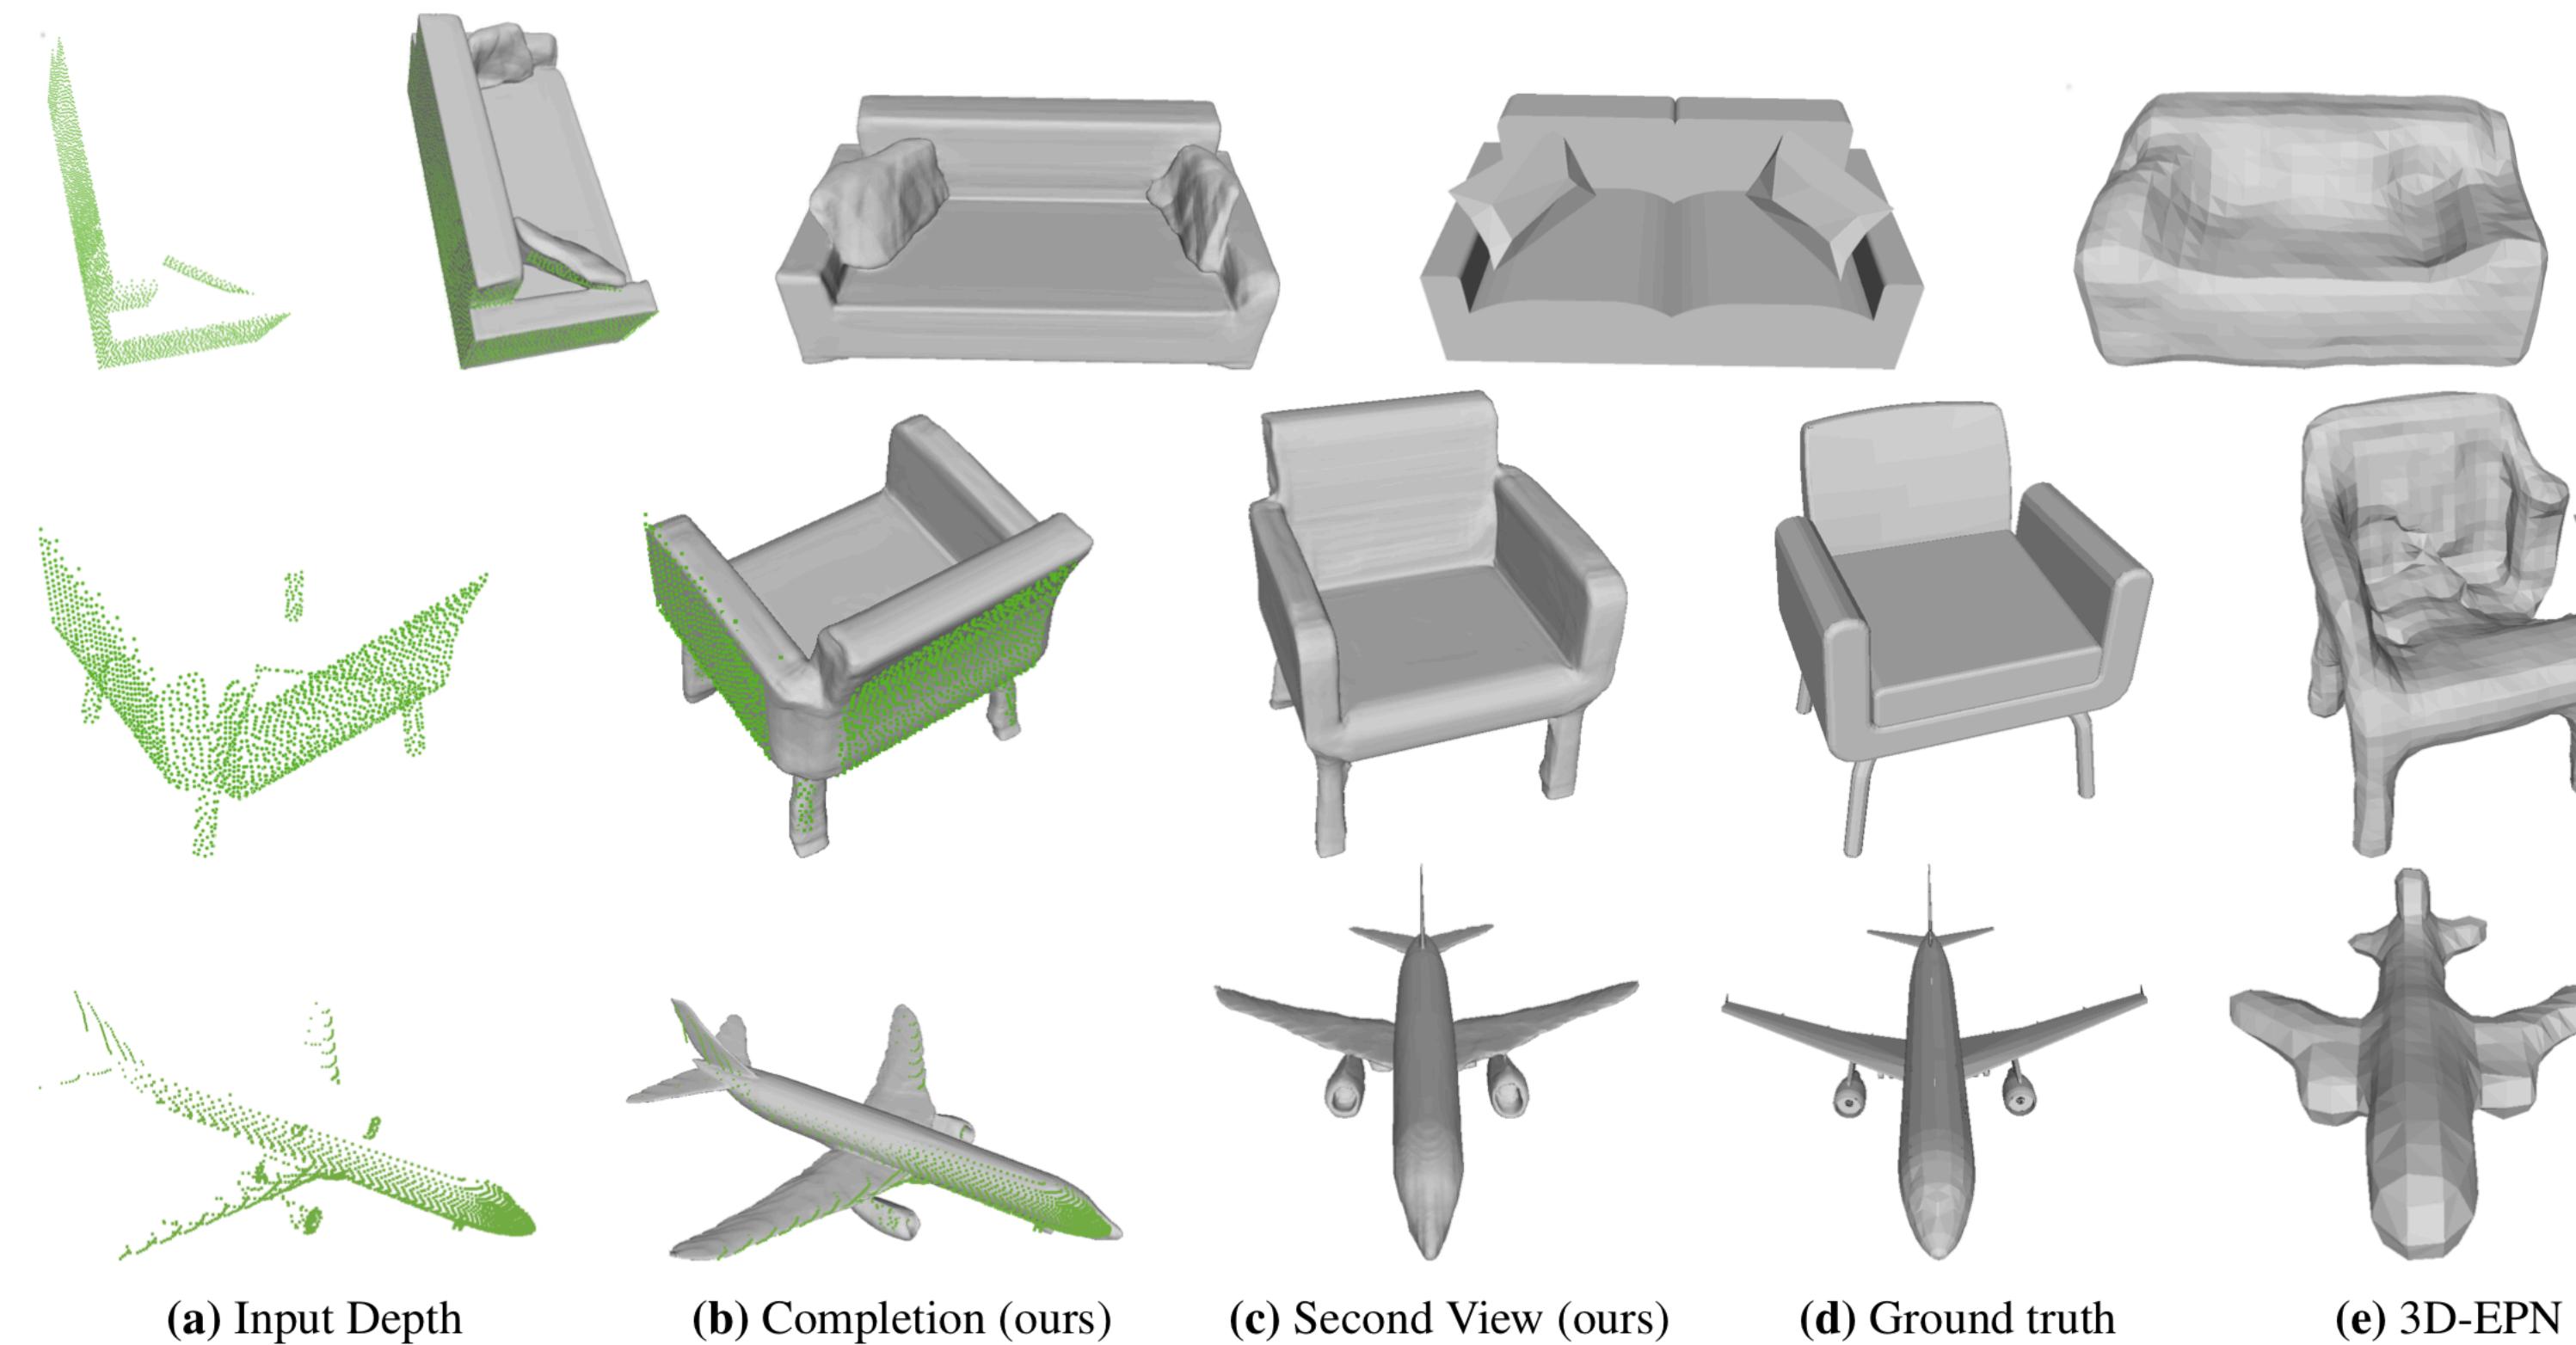
From the paper – shape reconstruction



**Figure 6:** Reconstruction comparison between DeepSDF and AtlasNet [22] (with 25-plane and sphere parameterization) for test shapes. Note that AtlasNet fails to capture the fine details of the chair, and that (f) shows holes on the surface of sofa and the plane.

# Expected results

From the paper – shape completion



**Figure 8:** For a given depth image visualized as a green point cloud, we show a comparison of shape completions from our DeepSDF approach against the true shape and 3D-EPN.

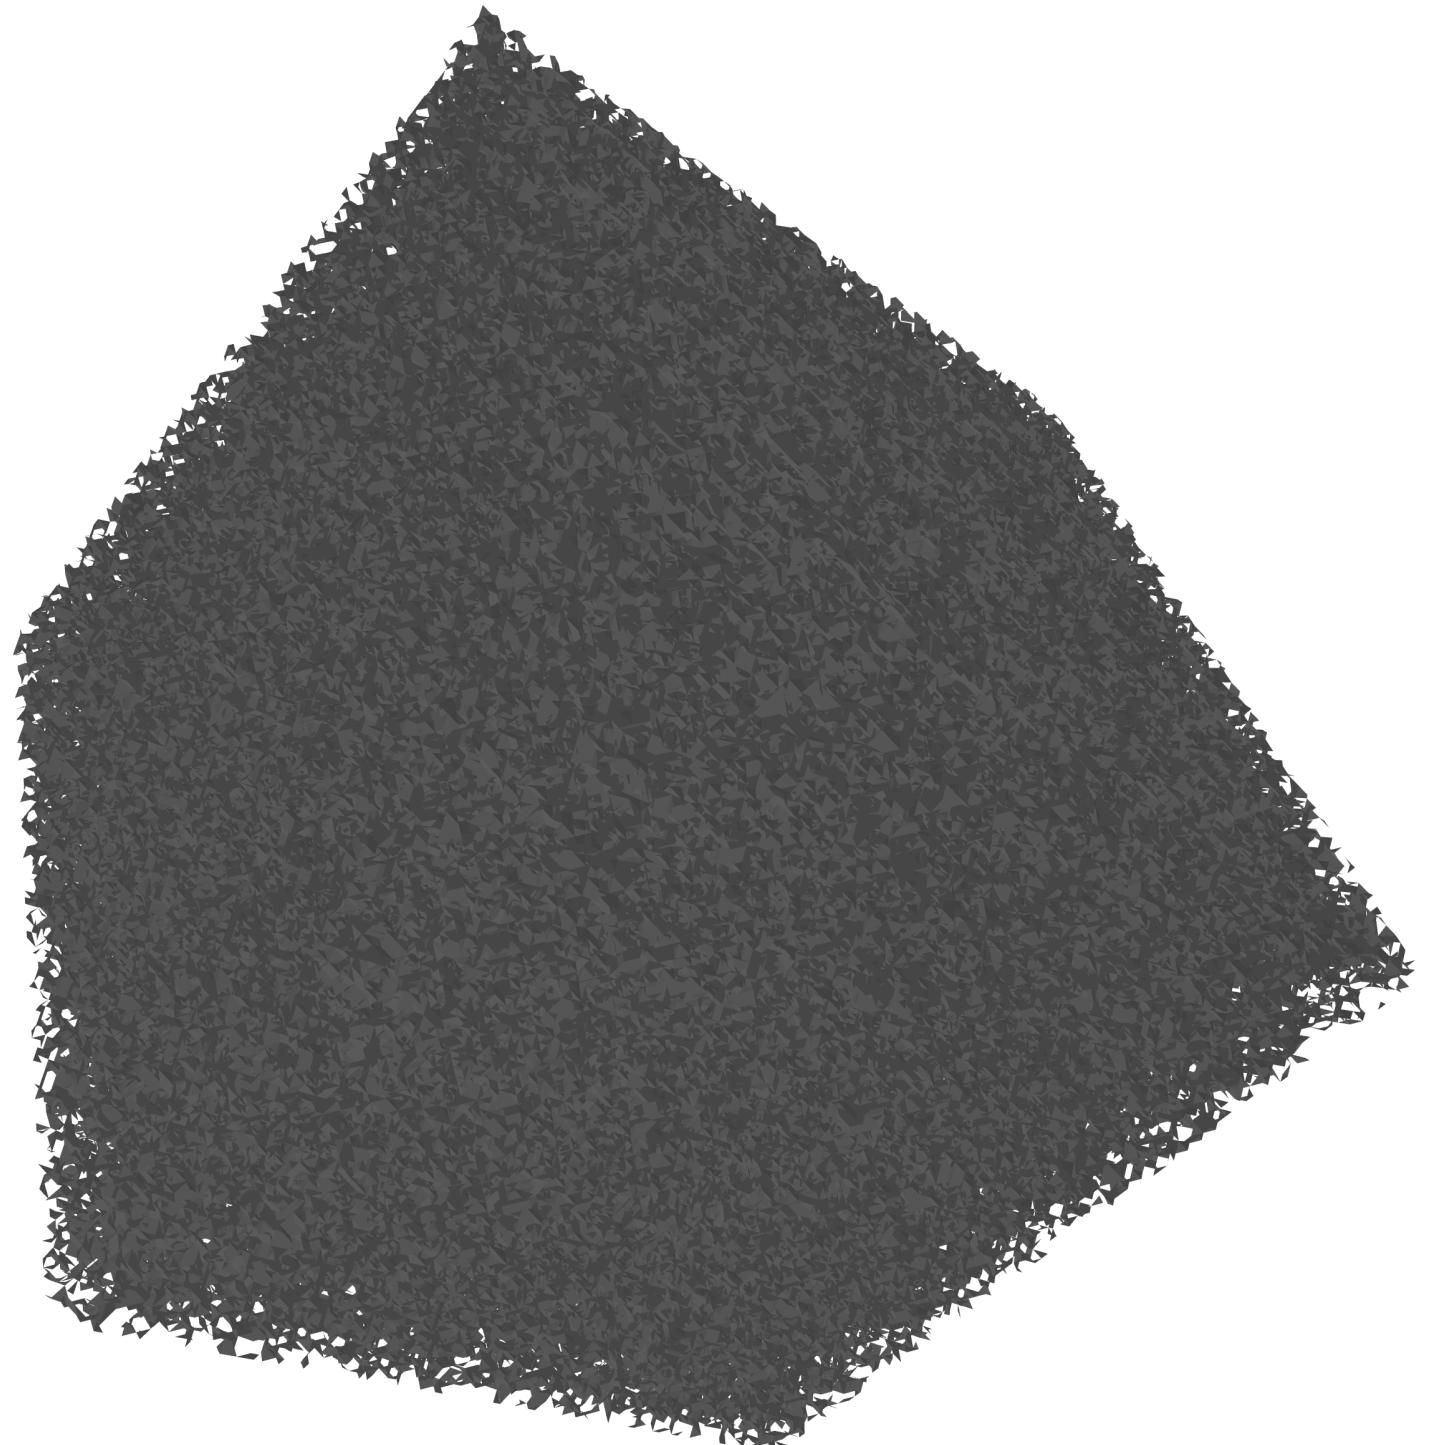
# Our results



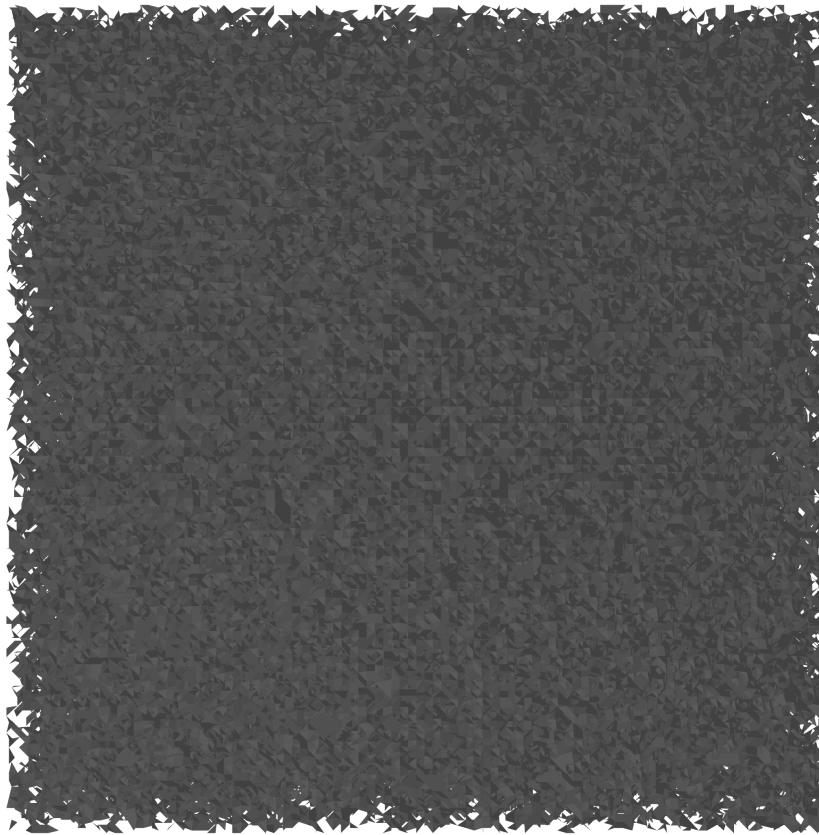
**Drum roll...**

# Our results

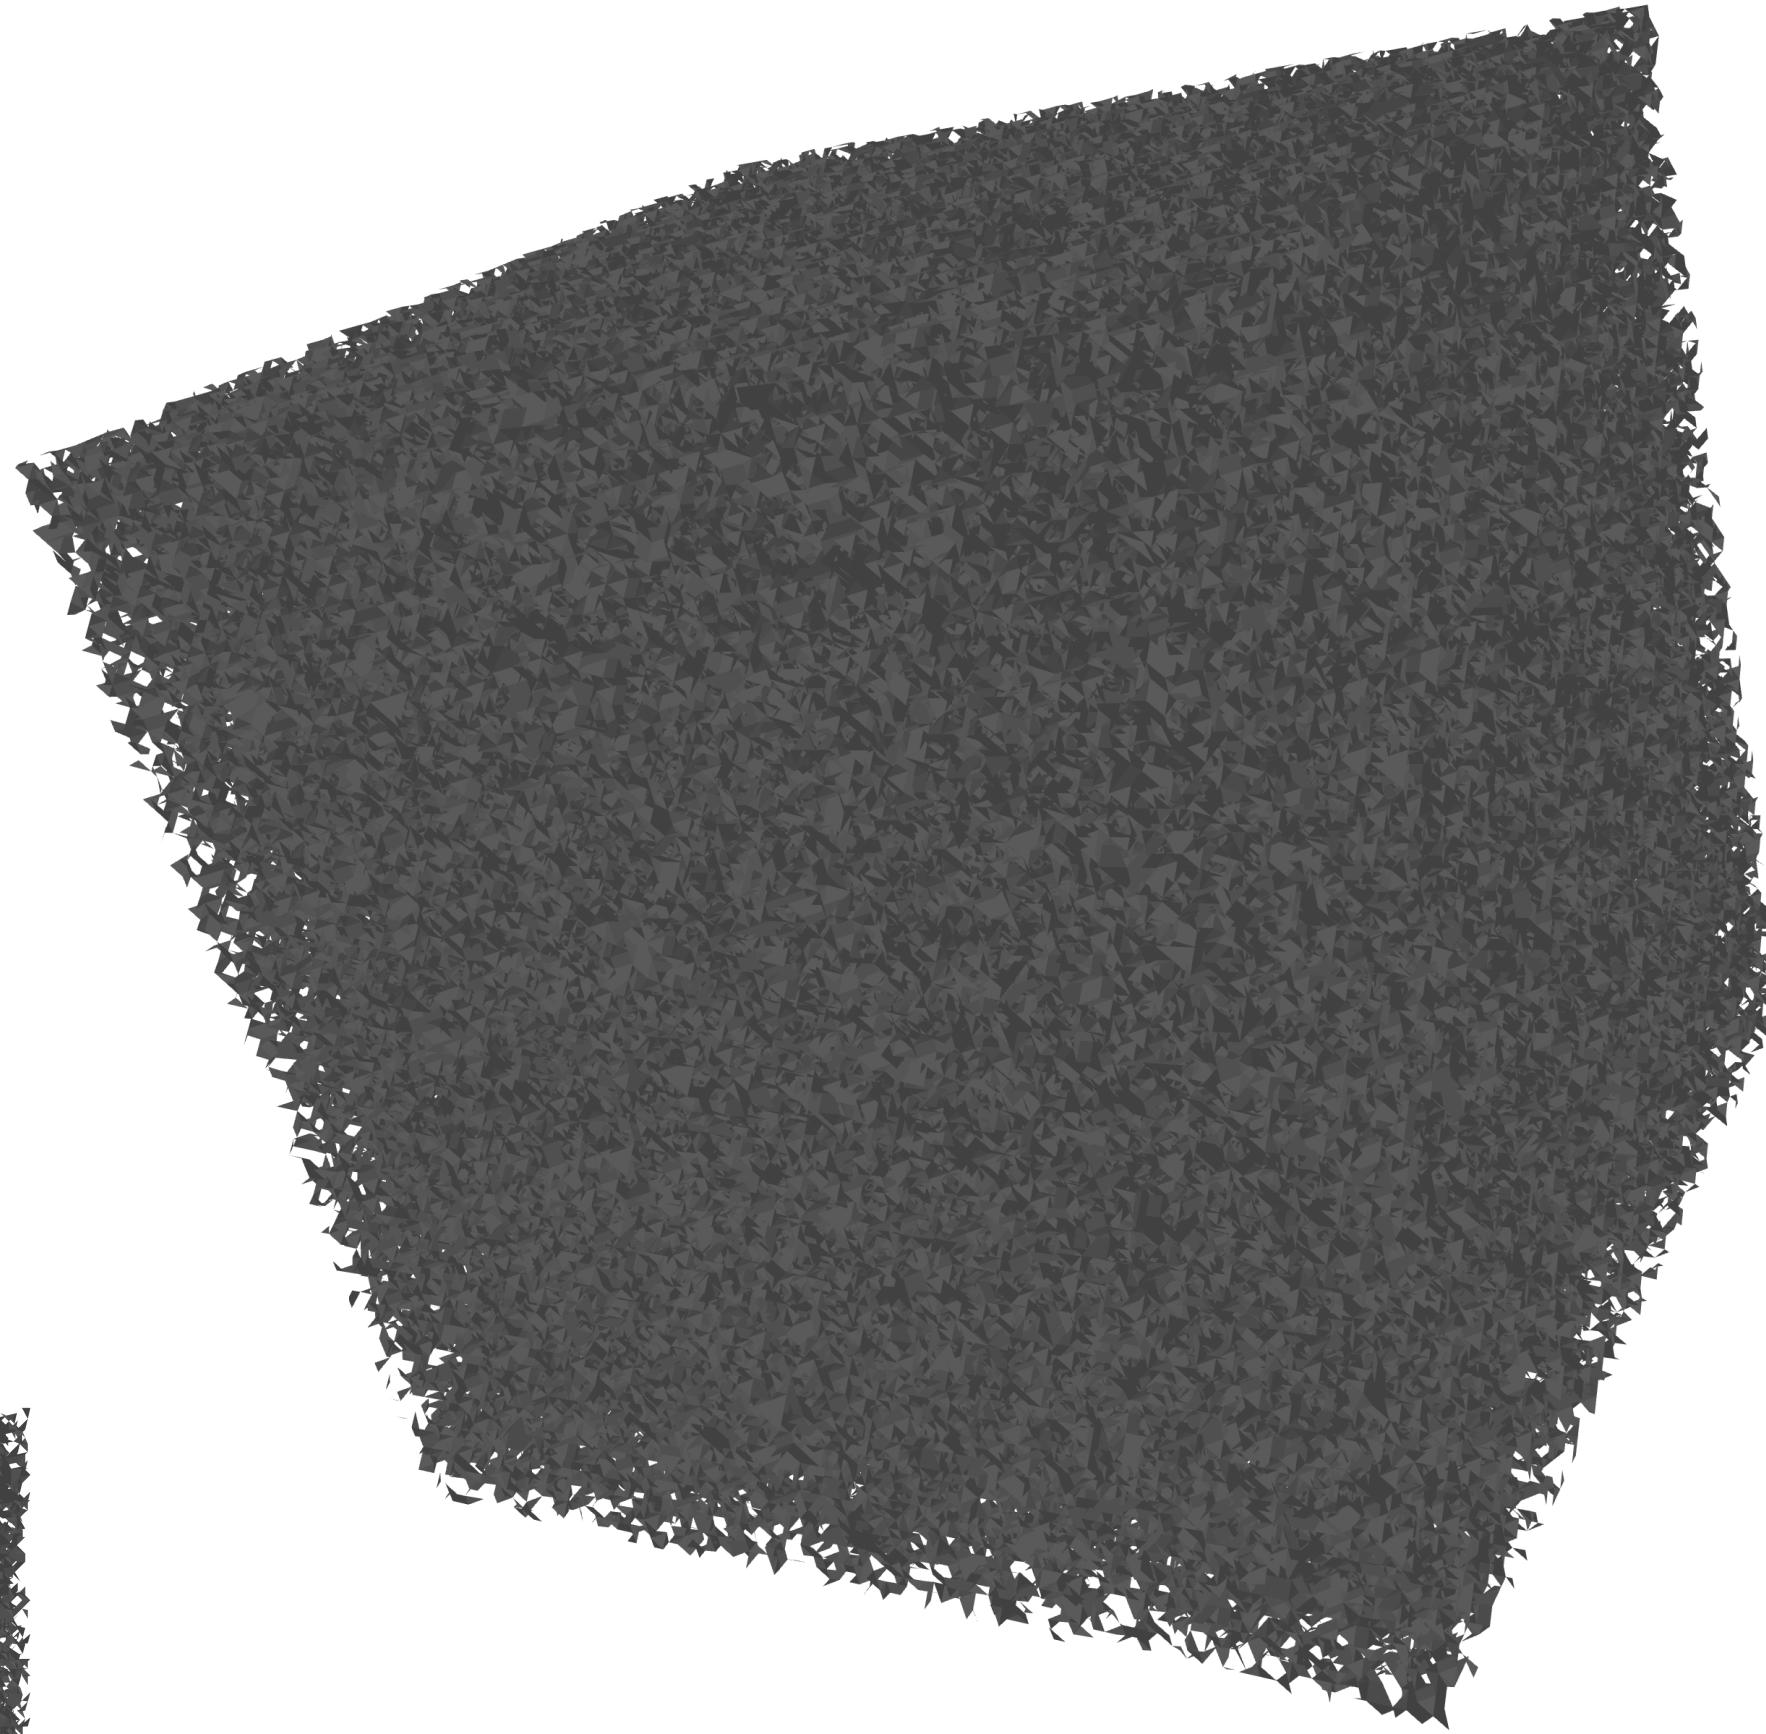
Not much... or maybe...



5 epochs, lr = 0.5



10 epochs, lr = 0.05



1000 epochs, lr = 0.005

# Our results

**Nope. This is an invisible horse.**

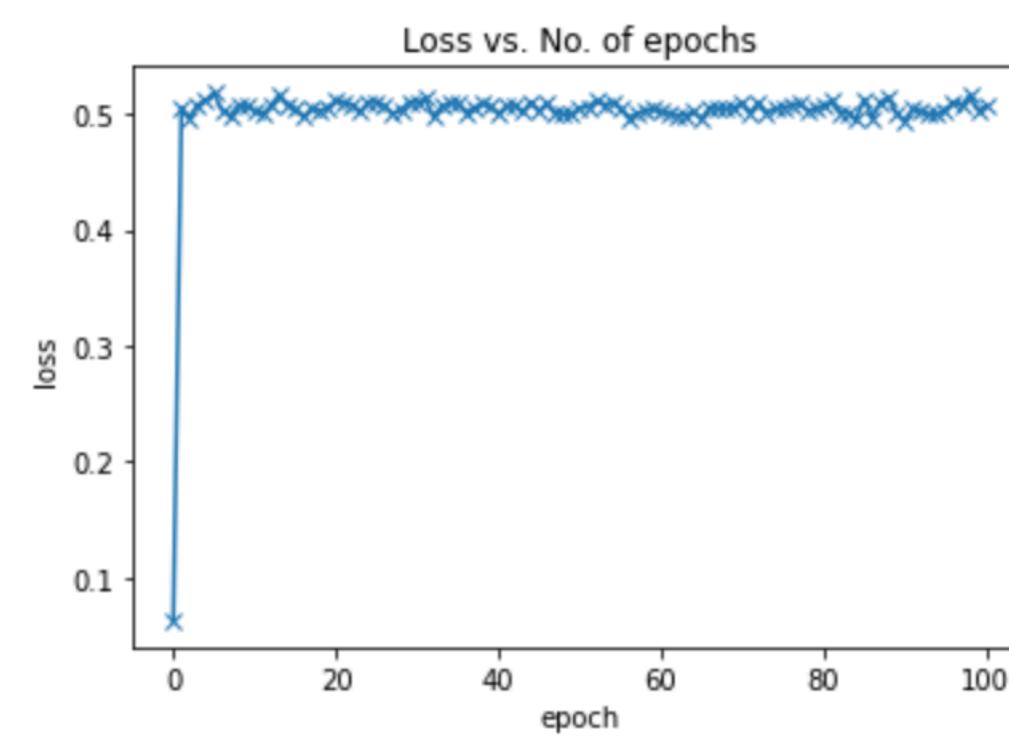
RuntimeError: No surface found at the given iso value.

10 epochs, lr = 0.005

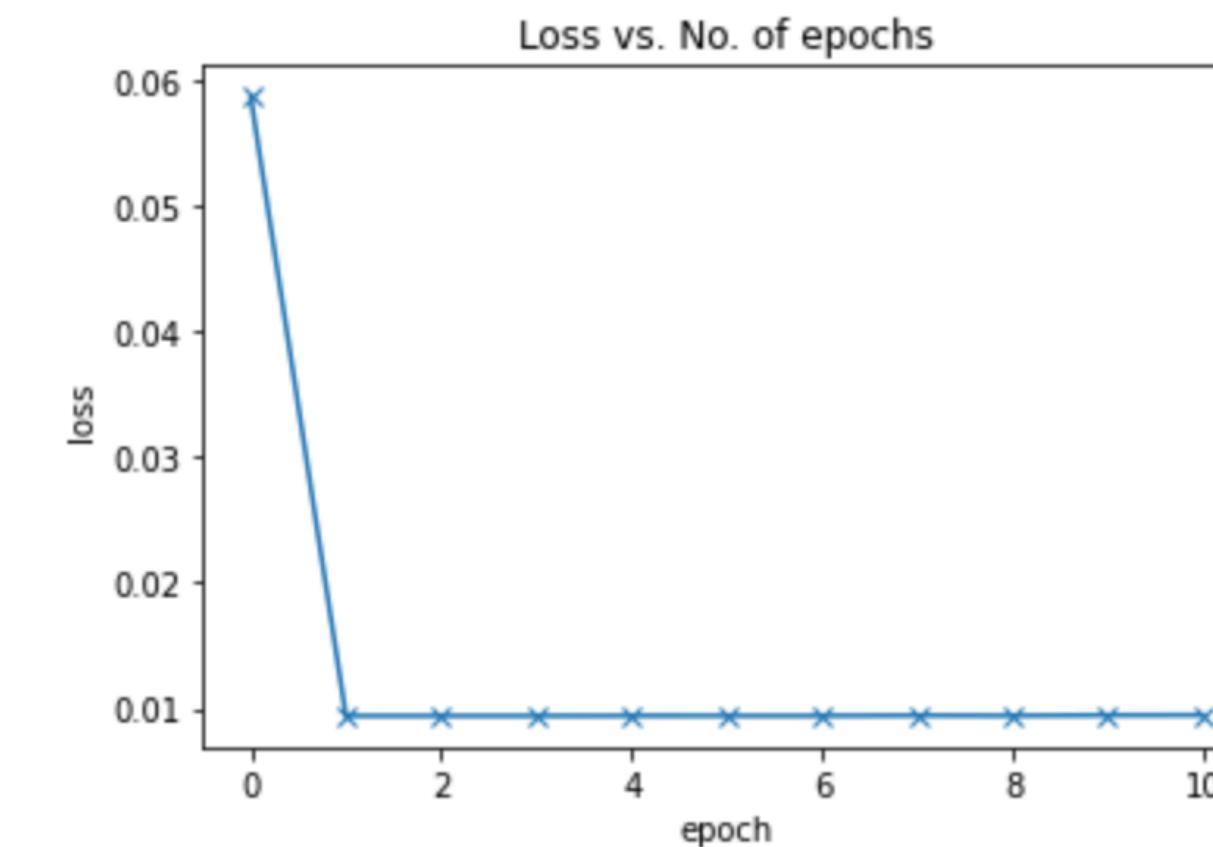
# Known issues

## Observations as of now

- For some reason, our loss does not seem to variate much
  - For the cube mesh, the loss varies between 0.48 and 0.52.
  - For the horse mesh, the loss just flatten at 0.01.



Training cube for 10 epochs

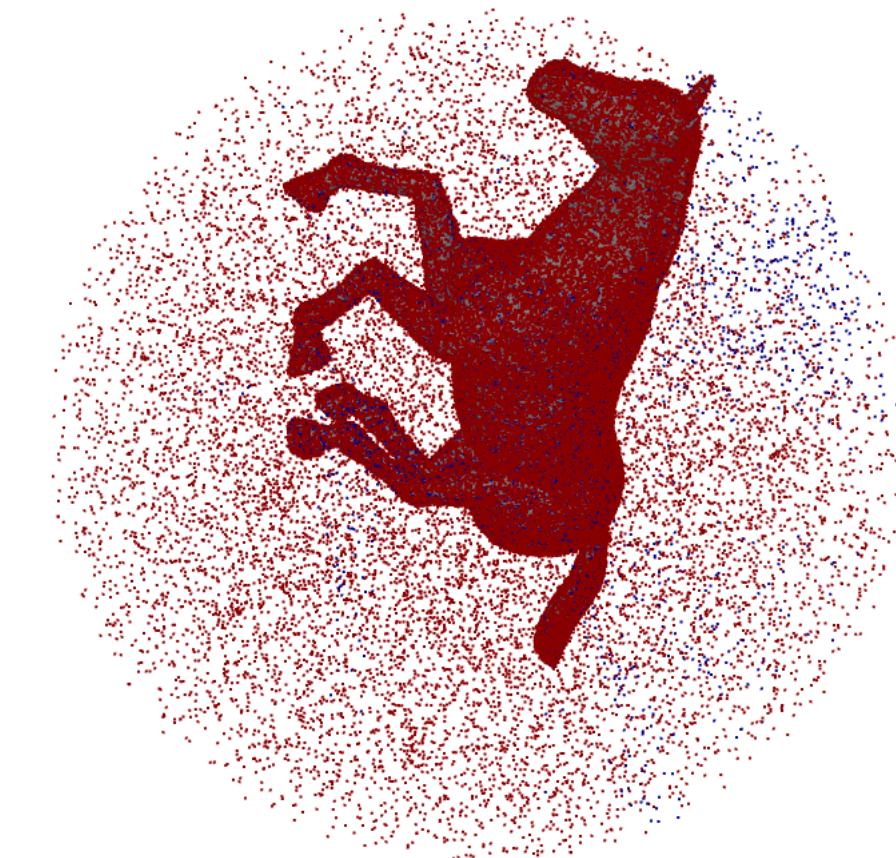


Training horse-1 for 10 epochs

# Known issues

## Observations as of now

- For the horse mesh, it seems that the SDF values for the spatial sample points are miscalculated.
- Some of the outer points are negative, whereas some inner points are positive.

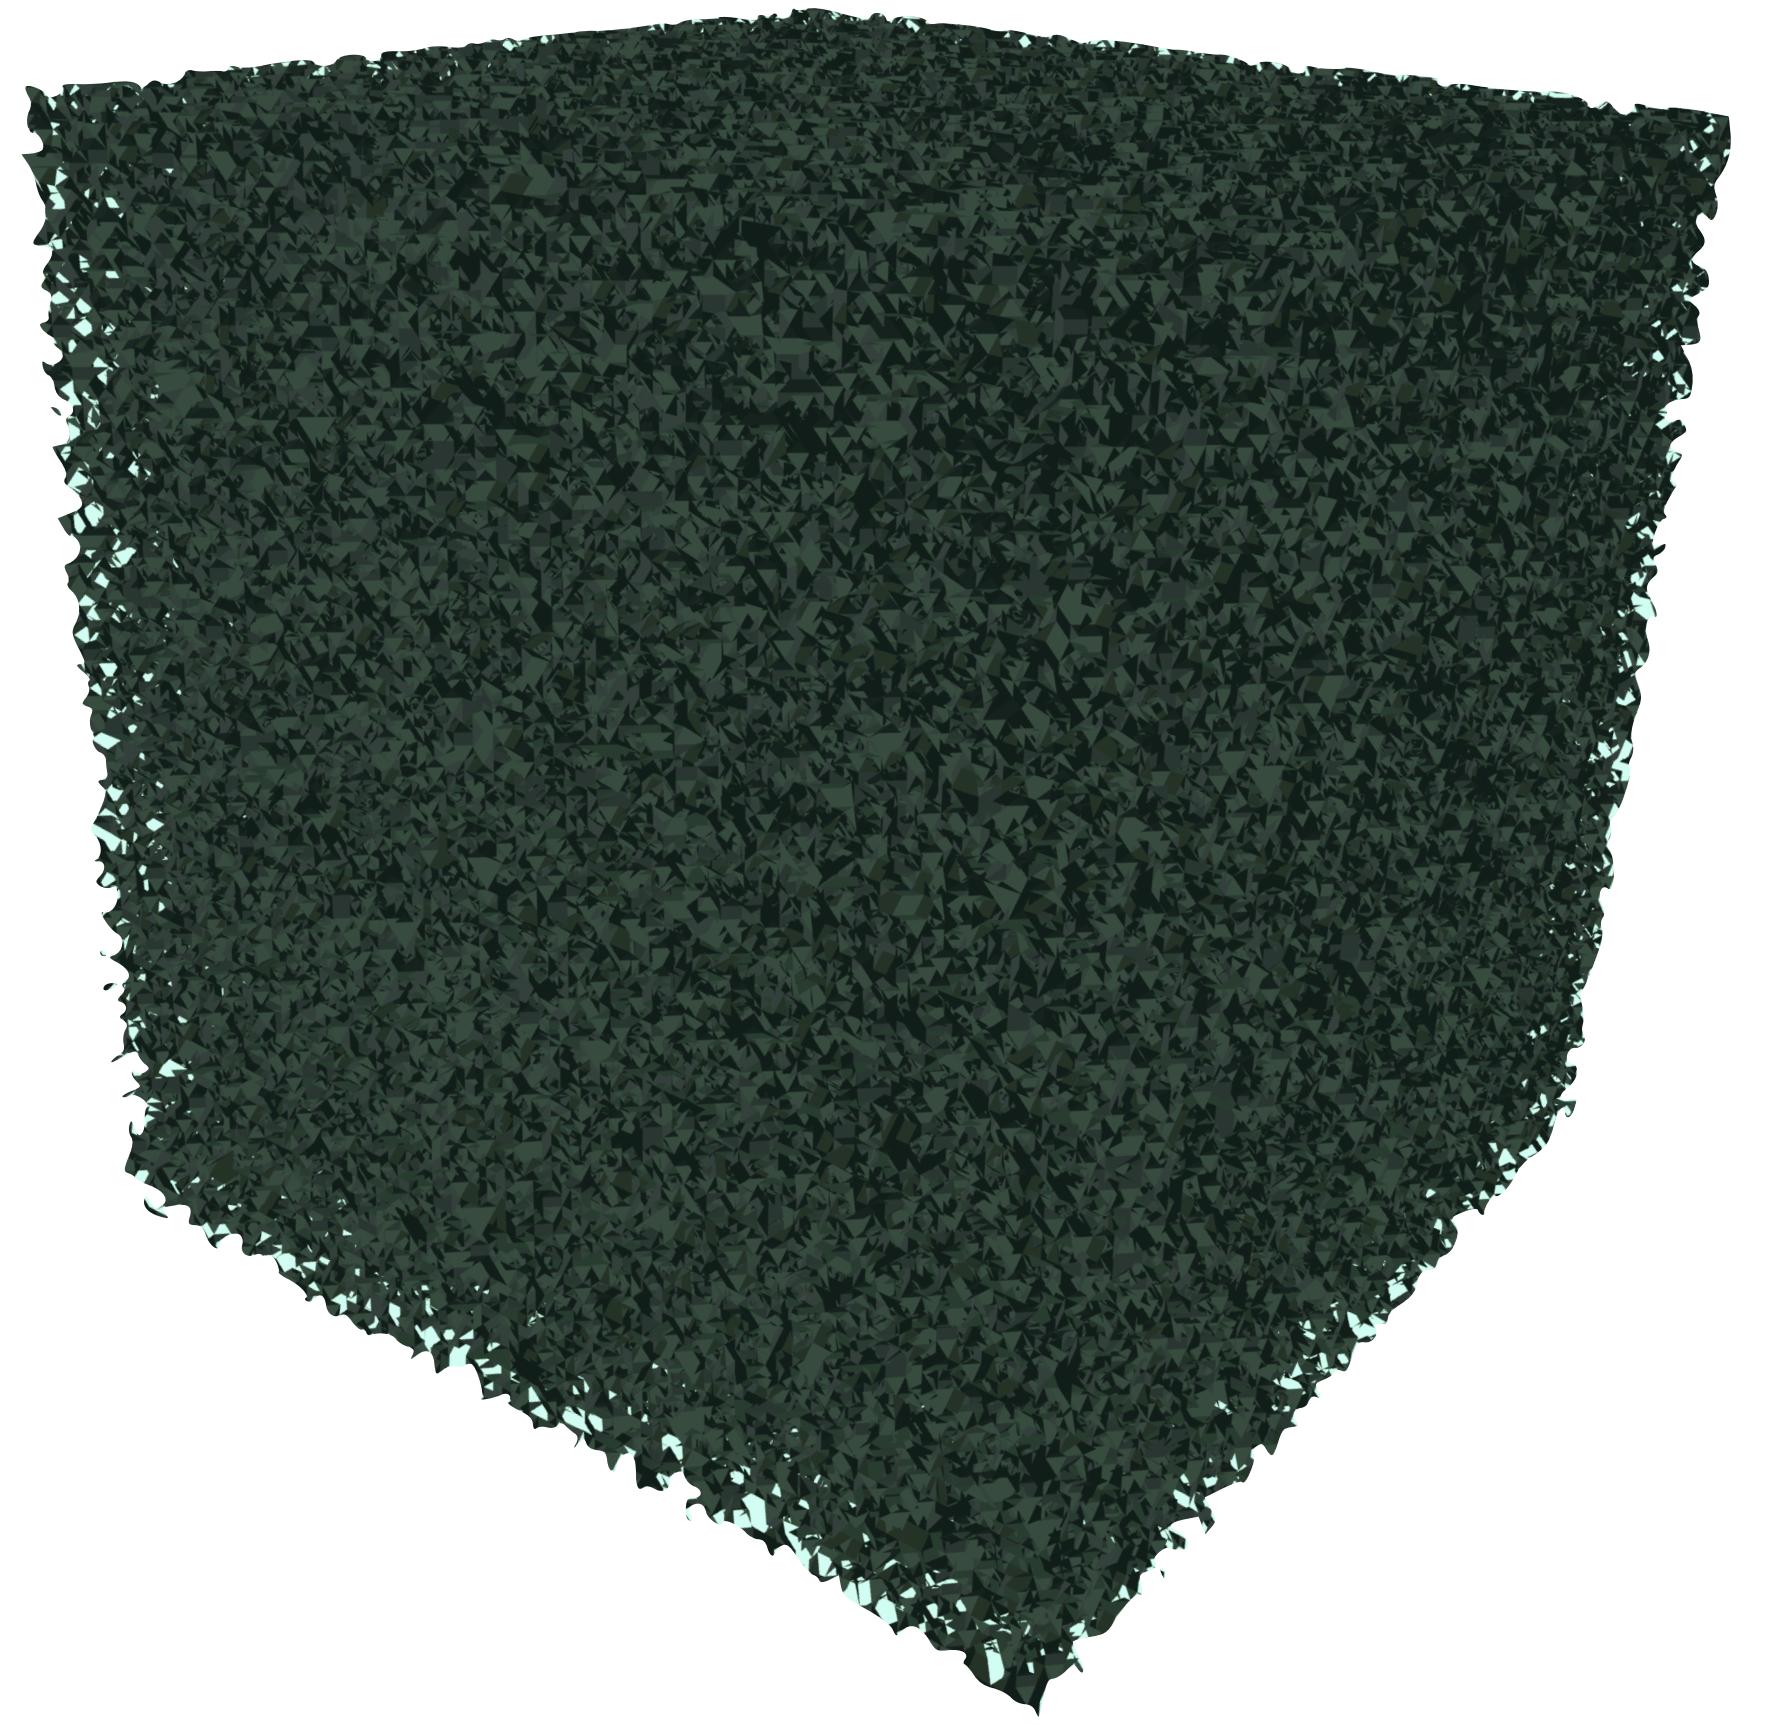


View of the spatial sample points for horse-1

# An idea of our extension

# Extension

- We did plan on experimenting with loss functions. Our goal was to maybe reduce training time or improve accuracy.
- However, much time has been spent debugging the existing loss function with still inconclusive results.
- Nevertheless, though beyond the scope of the original paper, it seems we may have found a slightly overengineered method to model some realistic cube bushes...



==>



**Thank you!**